

Masked Wheel Replacement Evaluation

This document describes the budgeted image-generation evaluation layer after the mask pipeline:

Roboflow / YOLO candidates

- > Qwen VLM keep/reject
- > final binary wheel mask
- > masked/reference-guided wheel replacement model

The evaluation runner is intentionally dry-run first. Paid generation only happens when `--execute` is passed.

Executive Summary

The practical question for this experiment was:

Given a car photo, a reference rim photo, and a Stage 2 wheel mask, which model can replace only the rims while preserving the car, tires, background, perspective, and low-light detail?

Important correction: the original 3-case frontier visual report is useful as a record of what happened, but it is **not a valid comparison of who best copied the silver target rim**. The manifest pointed at a silver reference image while `wheel_description` said the reference matte black multi-spoke wheel rim design. Text and image contradicted each other, so black-wheel outputs should not be treated as faithful silver-rim copying.

Corrected 3-case sweep status:

- Corrected manifest: `tmp/fal-inpaint-eval/ivan-corrected-silver-3cases.jsonl`
- Prompt now says: Use the exact wheel design, color, finish, spoke pattern, center cap, and material from the reference image
- studio lighting was removed; prompt now says Preserve original scene lighting
- Qwen is excluded from the honest main comparison because the current `fal-ai/qwen-image-edit/inpaint` config receives only `image_url`, `mask_url`, and `prompt`, not the reference wheel image.
- Direct Reve now has a `--mode production` eval path that sends exactly `[car, wheel]` and no mask, matching the current production Reve request shape.

Corrected short answer: **Reve production-equivalent is currently the best silver-rim visual copy in the 3-case corrected sweep, while Flux remains the strongest strict masked editor**. Flux keeps the edit localized through the mask, but still darkens the silver target on `ivan-C2` and `ivan-N2`.

Who Receives The Wheel Reference?




Figure 1: Corrected silver comparison

Model / path	Receives car	Receives mask	Receives wheel reference image	Main com- parison status
Flux Kontext via fal, flux-s085-g25	yes	yes	yes, reference_image	valid corrected masked compari- son
Reve direct API, --mode production	yes	no	yes, second reference_image item	valid production- equivalent compari- son
Reve direct API, masked mode Gemini via fal	yes	yes, as third reference image	yes	experimental, not production- equivalent
	yes	yes, as one of image_urls	yes	optional/secondary; previous run had time- out/instability
Qwen via fal, qwen-edit-default	yes	yes	no	excluded from honest wheel-copy compari- son

Model / path	Receives car	Receives mask	Receives wheel reference image	Main com- parison status
Z-Image / SDXL inpaint	yes	yes	no	prompt+mask baselines only

What was actually generated:

Family	Config / endpoint	Paid outputs?	Cases	Result
Flux Kontext	fal-ai/flux-kontext-lora/inpaint flux-s085-g25	yes	C1, C2, N2 plus wider sweeps	Best strict masked editor; corrected silver run completed
Flux tuning variants	flux-s075, s080, s085, guidance variants	yes	daylight + night sweeps	Useful for tuning; s075 helps night readability
Qwen Image Edit	fal-ai/qwen-image-edit/inpaint	yes	C1, C2, N2	Historical only; did not receive wheel reference image
Gemini image edit	fal-ai/gemini-3-pro-image-edit	yes	C1, C2, N2 timeout	Not usable in this masked setup
Reve direct API, masked	https://api.reve.com/v1/image/C1, C2, N2 with [car, wheel, mask]	yes	C1, C2, N2	Historical only; not production- equivalent
Reve direct API, production- equivalent	https://api.reve.com/v1/image/C1, C2, N2 with [car, wheel]	yes	C1, C2, N2	Best corrected silver-rim visual copy so far

Family	Config / endpoint	Paid outputs?	Cases	Result
Reve via fal.ai	<code>fal-ai/reve/fast/reim</code>	no, attempted	C1, C2, N2	Initial schema failed; config corrected afterward
Z-Image	<code>fal-ai/z-image/turbo/inpaint</code>	yes	wider comparison	Often drifts from black rim target
SDXL inpaint	<code>fal-ai/inpaint</code>	yes	model candidate sweep	Weak due to fram- ing/aspect changes
OpenAI image edit	<code>gpt-image-2 runner</code>	no paid outputs yet	dry-run only	Runner exists, but no generation examples yet

So to answer the coverage question directly: paid outputs exist for Flux, Qwen, Gemini, and Reve, but only Flux and production-equivalent Reve are in the current honest main comparison for copying the silver reference wheel. Qwen paid outputs are excluded because that endpoint did not receive the wheel image.

Visual TL;DR

The three comparison cases were:

- `ivan-C1` - white car, daylight.
- `ivan-C2` - dark gray SUV, daylight.
- `ivan-N2` - dark car, night.

Each case uses a car source image, a Qwen VLM-filtered binary mask, and the same rim reference image.

Reference rim used by the runs:

Historical frontier comparison from the flawed prompt/image mismatch run:

Zoomed wheel regions:

Reading the historical matrix:

- **Flux** edits the intended wheel areas most reliably.




Figure 2: Input cars, candidate overlays, and final VLM masks



Figure 3: Reference rim




Figure 4: Frontier comparison contact sheet

- **Qwen** keeps framing, but it did not receive the wheel reference image in this setup, so it cannot be judged as a faithful target-wheel copy.
- **Gemini** is unstable for this task: one daylight case became a two-image collage, and another inserted a standalone product wheel instead of editing the car.
- **Reve direct masked** preserves the car framing, but this path was experimental and not production-equivalent.

Inputs And Masks

The masks below are the exact Stage 2 VLM-filtered wheel masks used for the frontier comparison. White pixels indicate editable wheel areas; black pixels should be preserved.

ivan-C1






Figure 5: Frontier comparison zoom sheet



Figure 6: ivan-C1 VLM candidate overlay

ivan-C2





Figure 7: ivan-C2 VLM candidate overlay

ivan-N2



Generation Examples By Case

These rows show the source, mask, and generated outputs from the historical frontier comparison. They are kept for auditability, but the run had the silver-reference / matte-black-text mismatch described above. **Gemini** is missing on **ivan-N2** because that paid request timed out at 300 seconds.

ivan-C1

Visual notes:



Figure 8: ivan-N2 VLM candidate overlay



Figure 9: ivan-C1 generation row

- Flux produces the most direct masked replacement.
- Qwen changes the wheels, but the rim color/detail drifts away from the black reference.
- Gemini returned a two-image collage-like result, which is unusable for the product flow.
- Reve keeps the car image intact, but the wheels are closer to flat black disks than readable multi-spoke rims.

ivan-C2



Figure 10: ivan-C2 generation row

Visual notes:

- Flux is again the strongest masked replacement.
- Qwen makes the wheel structure too bright/greenish.
- Gemini inserts a standalone product-style wheel instead of editing the masked wheels on the car.

- Reve preserves the car and background, but the wheels are too dark and lose reference-rim detail.

ivan-N2



Figure 11: ivan-N2 generation row

Visual notes:

- Flux is dark, but still the best current default for the black-rim target.
- Qwen preserves more spoke brightness, but drifts strongly from the black rim target.
- Gemini timed out.
- Reve completed, but shows reddish wheel artifacts and still loses spoke readability.

Manifest

Create a JSONL manifest where each line is one case:

```
{"id":"case-001","car_image":"data/cars/001.jpg","mask_image":"tmp/masks/001.png","reference_image":"data/cars/001.jpg"}
```

Required fields:

- id
- car_image
- mask_image
- reference_image

Optional field:

- wheel_description - used in the text prompt. Keep this neutral unless it is machine-checked against target-wheel metadata. flux-kontext also receives the reference image directly; z-image uses prompt + mask only.

Prepare From Wheel-Labeling Dataset

For the local dataset:

/Users/nikolai/Downloads/wheel_labeling_nikolai_70/images

and wheel references:

/Users/nikolai/Documents/Dream Wheel AI/Ivan's Dataset/cars/rim1.jpg

build a budget-friendly 50-case manifest:

```
.venv/bin/python scripts/prepare_fal_eval_manifest.py \
  "/Users/nikolai/Downloads/wheel_labeling_nikolai_70/images" \
  "/Users/nikolai/Documents/Dream Wheel AI/Ivan's Dataset/cars/rim1.jpg" \
  --limit 50 \
  --max-long-edge 1024 \
  --output-dir tmp/fal-inpaint-eval/cases-rim1-1024 \
  --manifest tmp/fal-inpaint-eval/cases-rim1-1024.jsonl \
  --wheel-description "the exact wheel design, color, finish, spoke pattern, center cap, and"
```

This creates free proxy masks from XML wheel boxes. These masks are not the final Roboflow/Qwen masks, but the manifest format is identical, so generated masks can be swapped in later.

Review before inference:

```
tmp/fal-inpaint-eval/cases-rim1-1024/contact_sheet.jpg
```

Dry Run

```
.venv/bin/python scripts/fal_inpaint_eval.py cases.jsonl \
  --preset wide \
  --limit 50 \
  --max-estimated-cost 2.50
```

Outputs:

- tmp/fal-inpaint-eval/fal_inpaint_plan.jsonl
- tmp/fal-inpaint-eval/fal_inpaint_plan.csv

The script checks:

- mask image exists;
- mask size equals car image size;
- mask is not empty;
- mask does not cover more than 35% of the image;
- estimated cost is under --max-estimated-cost.

Current dry-run numbers for cases-rim1-1024.jsonl:

- wide, 50 cases, z-default + flux-s085-g25: 100 planned requests, 2 skipped by preflight, estimated cost \$1.7075.
- tuning, first 5 cases: 25 planned requests, estimated cost \$0.4177.

Paid Inference

```
FAL_KEY=... .venv/bin/python scripts/fal_inpaint_eval.py cases.jsonl \
  --preset wide \
  --limit 50 \
  --max-estimated-cost 2.50 \
  --execute
```

Results are written to:

- `tmp/fal-inpaint-eval/fal_inpaint_results.jsonl`
- `tmp/fal-inpaint-eval/fal_inpaint_results.csv`

Presets

wide:

- `z-default`
- `flux-s085-g25`

Use this for 50-case coverage.

flux-sweep:

- `flux-s075-g25`
- `flux-s080-g25`
- `flux-s085-g25`
- `flux-s080-g35`
- `flux-default`

Use this for a focused Flux parameter sweep on a small set of real VLM masks.

model-candidates:

- `flux-s085-g25`
- `flux-dev-s085-g35`
- `flux-general-s085-g35`
- `z-default`
- `sdxl-default`

Use this for a small cross-model smoke comparison. `flux-s085-g25` is the only candidate in this preset that receives the wheel reference image directly.

night-flux:

- `flux-s075-g25`
- `flux-s080-g25`
- `flux-s085-g25`
- `flux-s080-g35`
- `flux-s085-g35`

Use this for night-only tuning where wheel readability competes with black-rim style adherence.

frontier-edit:

- `flux-s085-g25`
- `qwen-edit-default`
- `reve-remix-rim-mask`
- `gemini-3-pro-rim-mask`

Use this only for a small comparison set. It includes expensive/slow image edit models and does not have the same strict mask semantics as Flux inpainting.

tuning:

- z-default
- z-strong
- flux-default
- flux-guidance-45
- flux-quality

Use this only on a small subset, usually 5 cases.

Budget Notes

The runner estimates cost from car image megapixels.

- fal-ai/flux-kontext-lora/inpaint: \$0.035/MP
- fal-ai/z-image/turbo/inpaint: \$0.01/MP
- fal-ai/qwen-image-edit/inpaint: \$0.03/MP
- fal-ai/reve/fast/remix: local planning estimate \$0.01/image
- fal-ai/gemini-3-pro-image-preview/edit: local planning estimate \$0.15/image

The Z-Image estimate is conservative because fal pages currently show both \$0.005/MP and \$0.01/MP in different places.

Flux Parameter Sweep Notes

The first paid Flux sweep used three real VLM-mask cases:

- ivan-C1 - white car, daylight.
- ivan-C2 - dark gray SUV, daylight.
- ivan-N2 - dark car, night.

Command shape:

```
.venv/bin/python scripts/fal_inpaint_eval.py \  
  tmp/fal-inpaint-eval/ivan-flux-param-tune-3cases.jsonl \  
  --preset flux-sweep \  
  --output-dir tmp/fal-inpaint-eval/results-flux-param-tune-real-masks \  
  --max-estimated-cost 0.40 \  
  --execute
```

Observed outputs:

- 15/15 requests completed.
- Total estimated cost: \$0.3021.
- Local comparison sheets:
 - tmp/fal-inpaint-eval/results-flux-param-tune-real-masks/flux_param_tune_contact_sheet
 - tmp/fal-inpaint-eval/results-flux-param-tune-real-masks/flux_param_tune_zoom_sheet

Current visual read:

- `flux-default` is a useful baseline, but not the best production candidate. It tends to make night wheels too dark/empty.
- `flux-s080-g25` preserves more bright spoke detail, but often drifts away from the black reference rim.
- `flux-s080-g35` and `flux-default` are darker and more prompt-adherent, but can lose rim structure in low light.
- `flux-s085-g25` is the best current compromise across the three cases: black rim style, enough spoke structure, and fewer night-case failures than `flux-default`.

Important integration caveat: `fal-ai/flux-kontext-lora/inpaint` may return a different output resolution than the input image. In the first sweep, a 1024x594 input produced 1328x800 outputs. Any production integration should either request/verify exact sizing if the endpoint supports it or resize the final image back to the original dimensions before handing it to downstream code.

Model Candidate Sweep Notes

The first cross-model sweep used the same three real VLM-mask cases and initially tested:

- `flux-s085-g25`
- `flux-dev-s085-g35`
- `flux-krea-s085-g35`
- `flux-general-s085-g35`
- `z-default`
- `sdxl-default`

Observed status:

- 15/18 requests completed.
- `flux-krea-s085-g35` timed out on all three cases, so it was removed from the `model-candidates` preset.
- Current no-Krea dry-run: 3 cases, 15 planned requests, estimated \$0.2331.
- Local comparison sheet:
 - `tmp/fal-inpaint-eval/results-model-candidates-3cases/model_candidates_contact_sheet`

Early visual read:

- `flux-s085-g25` remains the best reference-driven candidate.
- `flux-dev-s085-g35` and `flux-general-s085-g35` are viable prompt+mask baselines, but they do not receive the wheel reference image.
- `z-default` often preserves/creates brighter spokes and can drift from the black rim target.
- `sdxl-default` changes framing/aspect enough that it is a weak production candidate without additional sizing controls.

OpenAI Image Edit Baseline

OpenAI image edits are tracked as a separate managed baseline because the API returns base64 image data and uses a different mask convention than fal.ai.

Runner:

```
.venv/bin/python scripts/openai_image_edit_eval.py \  
  tmp/fal-inpaint-eval/ivan-cars-rim1-1024-vlm-masks-selected.jsonl \  
  --limit 6 \  
  --output-dir tmp/openai-image-edit-eval/ivan-vlm-6cases
```

Paid run:

```
OPENAI_API_KEY=... .venv/bin/python scripts/openai_image_edit_eval.py \  
  tmp/fal-inpaint-eval/ivan-cars-rim1-1024-vlm-masks-selected.jsonl \  
  --limit 6 \  
  --output-dir tmp/openai-image-edit-eval/ivan-vlm-6cases \  
  --execute
```

The runner sends:

- first input image: the car image to edit;
- second input image: the rim reference image;
- mask: an RGBA alpha mask derived from the Stage 2 binary wheel mask.

Mask convention: Stage 2 masks are white where wheels should be edited. OpenAI image edits use transparent alpha pixels as editable regions, so the runner converts white wheel pixels to transparent and keeps all other pixels opaque.

Outputs:

- openai_image_edit_plan.jsonl
- openai_image_edit_plan.csv
- openai_masks/*.openai-alpha-mask.png
- openai_image_edit_results.jsonl
- openai_image_edit_results.csv
- outputs/*.png

Current local status: runner and dry-run path are implemented; paid execution requires OPENAI_API_KEY. Default model is gpt-image-2.

Night Flux Sweep Notes

The night-only Flux sweep used ivan-N1, ivan-N2, and ivan-N3 with real VLM masks:

```
.venv/bin/python scripts/fal_inpaint_eval.py \  
  tmp/fal-inpaint-eval/ivan-vlm-night-3cases.jsonl \  
  --preset night-flux \  
  --output-dir tmp/fal-inpaint-eval/results-flux-night-sweep \  
  --limit 6
```

```
--max-estimated-cost 0.40 \  
--execute
```

Observed outputs:

- 15/15 requests completed.
- Total estimated cost: \$0.3097.
- Local comparison sheets:
 - tmp/fal-inpaint-eval/results-flux-night-sweep/night_sweep_contact_sheet.jpg
 - tmp/fal-inpaint-eval/results-flux-night-sweep/night_sweep_zoom_sheet.jpg

Current visual read:

- flux-s075-g25 is the best night readability candidate. It keeps more spoke detail visible, especially on *ivan-N2*, where stronger settings can collapse into a black disk.
- flux-s085-g25 remains the better all-around/default candidate because it adheres more strongly to the black reference rim style on daylight cases.
- flux-s080-g35 can improve contrast on some night wheels, but it is less stable and produced an obviously wrong steel-wheel-like result on *ivan-N1*.
- flux-s085-g35 is usually too dark for night cases.

Reference Candidate Sweep Notes

The short reference-conditioning sweep used the same three cases as the model candidate sweep:

- *ivan-C1* - white car, daylight.
- *ivan-C2* - dark gray SUV, daylight.
- *ivan-N2* - dark car, night.

It compared the current Flux Kontext candidates against two fal.ai reference alternatives:

- flux-general-reference-rim - fal-ai/flux-general/inpainting with reference_image_url.
- qwen-edit-default - fal-ai/qwen-image-edit/inpaint.

Command:

```
.venv/bin/python scripts/fal_inpaint_eval.py \  
tmp/fal-inpaint-eval/ivan-flux-param-tune-3cases.jsonl \  
--preset reference-candidates \  
--output-dir tmp/fal-inpaint-eval/results-reference-candidates-3cases \  
--max-estimated-cost 0.40 \  
--execute \  
--client-timeout 240
```

Observed outputs:

- 12/12 requests completed.
- Total estimated cost: \$0.2331.
- Local comparison sheets:
 - tmp/fal-inpaint-eval/results-reference-candidates-3cases/reference_candidates_conta
 - tmp/fal-inpaint-eval/results-reference-candidates-3cases/reference_candidates_zoom

Current visual read:

- `qwen-edit-default` is fast and inexpensive, but it appears weaker for faithful rim replacement in this masked workflow.
- `flux-general-reference-rim` accepts the built-in reference parameter and is a useful experiment, but its output resolution differs from the input and still needs stricter visual review before it can replace Flux Kontext.
- `flux-s085-g25` remains the safest all-around production candidate.
- `flux-s075-g25` remains the best night readability fallback.

Frontier Edit Sweep Notes

The frontier edit sweep compared the current Flux/Qwen candidates against Gemini and direct Reve remix on the same three cases:

- `ivan-C1` - white car, daylight.
- `ivan-C2` - dark gray SUV, daylight.
- `ivan-N2` - dark car, night.

fal.ai command:

```
.venv/bin/python scripts/fal_inpaint_eval.py \
  tmp/fal-inpaint-eval/ivan-flux-param-tune-3cases.jsonl \
  --preset frontier-edit \
  --output-dir tmp/fal-inpaint-eval/results-frontier-edit-3cases \
  --max-estimated-cost 0.80 \
  --execute \
  --client-timeout 300
```

Direct Reve command:

```
.venv/bin/python scripts/reve_image_edit_eval.py \
  tmp/fal-inpaint-eval/ivan-flux-param-tune-3cases.jsonl \
  --limit 3 \
  --output-dir tmp/reve-image-edit-eval/ivan-vlm-3cases-direct \
  --execute \
  --timeout 240
```

Observed outputs:

- fal frontier sweep: 8/12 completed.
- `flux-s085-g25`: 3/3 completed.
- `qwen-edit-default`: 3/3 completed.

- **gemini-3-pro-rim-mask**: 2/3 completed; **ivan-N2** timed out at 300 seconds.
- **reve-remix-rim-mask** through fal initially failed validation because the endpoint expected `image_urls`; the local config was corrected afterward.
- direct Reve: 3/3 completed after retrying one transient SSL failure.
- Local comparison sheets:
 - `tmp/reve-image-edit-eval/frontier-comparison-3cases/frontier_comparison_contact_sheet`
 - `tmp/reve-image-edit-eval/frontier-comparison-3cases/frontier_comparison_zoom_sheet`

Historical visual read, before the silver-reference prompt mismatch was found:

- **flux-s085-g25** remains the best all-around candidate in this comparison.
- **qwen-edit-default** keeps the car framing, but tends to make rim structure too bright/greenish. This is not a valid target-wheel-copy judgment because Qwen did not receive the wheel reference image.
- **gemini-3-pro-rim-mask** is not usable as a masked wheel-edit baseline in this setup: one output became a two-image collage, and another inserted a standalone product-style wheel instead of editing only the masked rims.
- Direct Reve technically works and preserves the car framing, but the outputs from this masked experimental path are not production-equivalent.

Corrected Silver Sweep Notes

This is the corrected follow-up sweep for the silver `rim1.jpg` reference. It removes the matte-black text mismatch and uses a neutral wheel description:

the exact wheel design, color, finish, spoke pattern, center cap, and material from the reference image.

Adapted universal prompt principle:

Replace the existing wheels on the car with the provided alloy wheel design.

Use the exact wheel design, color, finish, spoke pattern, center cap, and material from the reference image. Match correct perspective, scale, and alignment with the car hub.

Preserve original scene lighting, realistic reflections, and shadows.

Keep brake disc and wheel depth physically plausible.

Match tire profile and maintain correct wheel size relative to the car.

Do not alter car body, color, or background.

Photorealistic, ultra detailed, automotive photography, high resolution, natural reflections

Corrected manifest:

```
tmp/fal-inpaint-eval/ivan-corrected-silver-3cases.jsonl
```

Flux corrected command:

```
.venv/bin/python scripts/fal_inpaint_eval.py \
  tmp/fal-inpaint-eval/ivan-corrected-silver-3cases.jsonl \
  --config flux-s085-g25 \
```

```
--output-dir tmp/fal-inpaint-eval/results-corrected-silver-flux \
--max-estimated-cost 0.10 \
--execute
```

Reve production-equivalent corrected command:

```
.venv/bin/python scripts/reve_image_edit_eval.py \
tmp/fal-inpaint-eval/ivan-corrected-silver-3cases.jsonl \
--mode production \
--output-dir tmp/reve-image-edit-eval/results-corrected-silver-production \
--execute
```

Observed status:

- Flux corrected: 3/3 completed, estimated cost \$0.0604.
- Reve production-equivalent corrected: 3/3 completed.
- Visual sheet:
 - docs/assets/fal-mask-inpaint-eval/corrected/corrected-silver-flux-vs-reve-production

Current visual read:

- Reve production-equivalent copies the silver target wheel color and spoke pattern more directly across the three corrected cases.
- Flux remains the better strict masked-edit candidate, but it still darkens the silver wheel on **ivan-C2** and especially **ivan-N2**.
- The fastest path to a good demo today is to show the corrected Reve production-equivalent result as the best visual silver-wheel transfer, while keeping Flux as the safer masked/inpaint control.
- Qwen should stay out of the main result unless a different Qwen endpoint is chosen that can receive the car image and wheel reference image. The available multi-image Qwen edit endpoint can be tested separately, but without a hard mask it would be another experimental path rather than the same comparison.